Image Classification of Letters

- A Model Comparison Approach -

•••

General Assembly Final Project 2018

by Mundy Otto Reimer

notMNIST Dataset

- 10 classes for letters A - J
- Different Glyphs of Fonts
- Size 28x28 pixels
- 19k hand-cleaned instances
- 500k uncleaned instances

Objective:

Classification Problem

.png Image -> Unicode Value

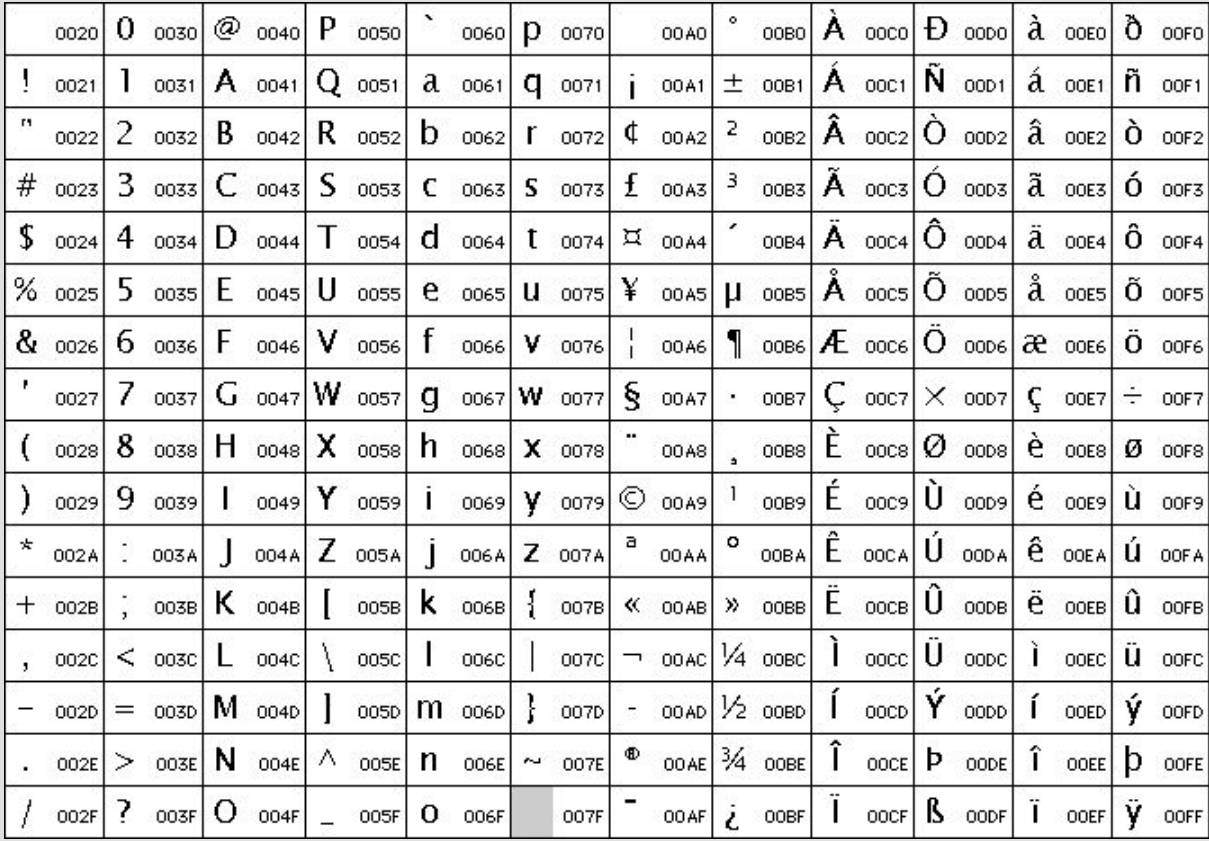

Cleaning the Data

Loading

- Un-Pickling
- Each class loaded into separate dataset (memory issues)
- Merged at end in 1 big data set

Normalizing

- Convert Dataset into 3D array

  (image index, x, y)
  Floating Point #s

- Normalized to ~zero mean & std ~0.5 (easier training)

Sampling Labeled Data

- Verify data & labels still ok

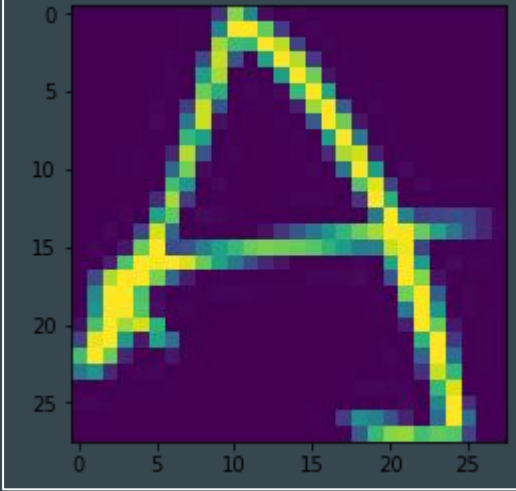

Partition & Process Data: Training > Validation > Test

Class Distribution

- Check data if balanced across letter classes (Accuracy Paradox)

52k per Large Letter

2k per Small Letter

Randomization

- Shuffle labels so training and test distributions match
- Plot each histogram

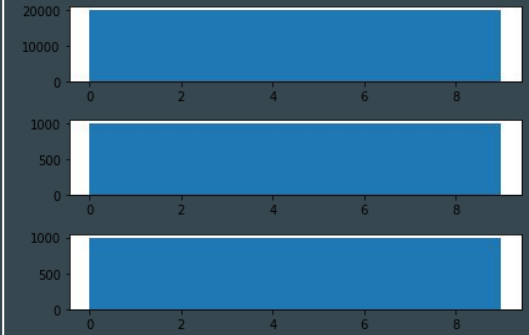

Measure Overlap

- After split of 90/5/5, check for Training, Validation, Test sample overlap

Duplicates in training set ~12k

Duplicates in train+val ~900

Duplicates in train+test ~1k

- Expect to use in environment w/ no overlap?

Model Comparisons

	Training / Val / Test Scores	
Logistic Regression	.76 / .80	• Cheap & Simple • Scores high (90s) for classic MNIST dataset
LR w/ Gradient Descent	.78 / .75/ .82	• Used softmax at end (R -> [0,1]) for multiple classes
LR w/ Stochastic Gradient Descent	.77 / .78 / .85	• 4.76 sec (as opposed to 15+ sec) • Little improvement in accuracy
1-Layer Feedforward Neural Network	.83 / .82/ .88	• Small accuracy improvement over faster & simpler LR w/ Stochastic GD • Computationally Expensive AND opaque

Model Comparisons

	Training / Val / Test Scores	
Neural Network w/ L2 Regularization	.86 / .85 / .91	• L2 Ridge (Used b/c large weights indicate possible overfitting?), so multiply them by small fraction (adds penalty on norm of the weights to loss)
Neural Dropout	.84 / .85 / .91	• Since improvement from overfitting, try another overfitting technique (didn't improve much) • Also used on fully-connected networks
Multiple Layers & Learning Rate Decay	.86 / .85 / .91	• Multiple layers too computationally expensive (+hours) • Instead implemented LRD to reduce training time back to ~10 sec
2+1-Layer Convolutional Neural Network	.69? / .86 / .93	• Computation time went back up (40 min) • Stride 2 to reduce dimensionality (13 min)

Convolutional Neural Networks

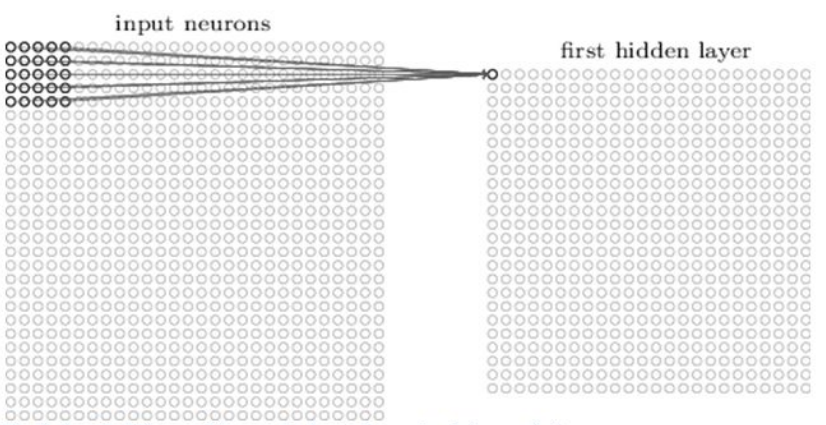

Visualization of 5 x 5 filter convolving around an input volume and producing an activation map


Original image

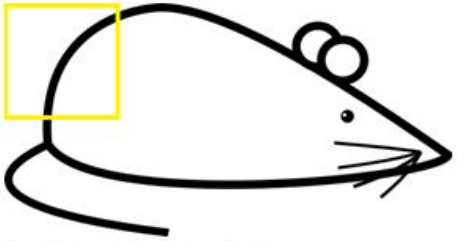
Visualization of the filter on the image

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter

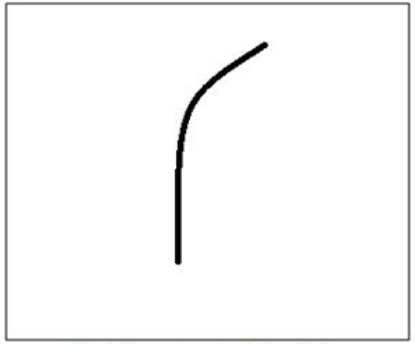
Visualization of a curve detector filter


Visualization of the receptive field

0	0	0	0	0	0	30
0	0	0	0	50	50	50
0	0	0	20	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0

Pixel representation of the receptive field

*

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter

Multiplication and Summation = (50*30)+(50*30)+(50*30)+(20*30)+(50*30) = 6600 (A large number!)

CNN cont: Stride & Padding

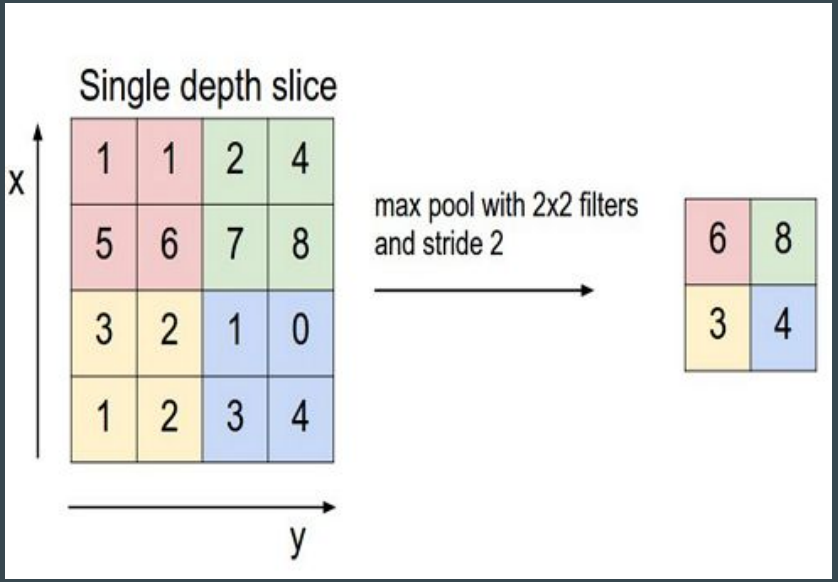

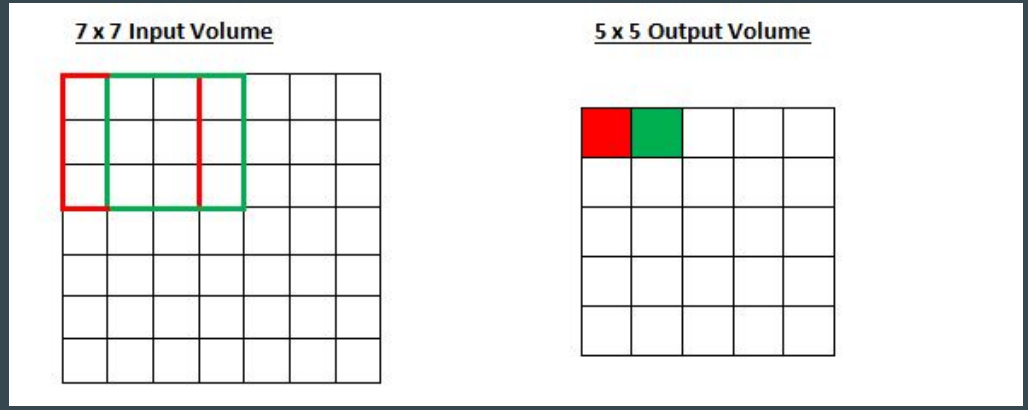

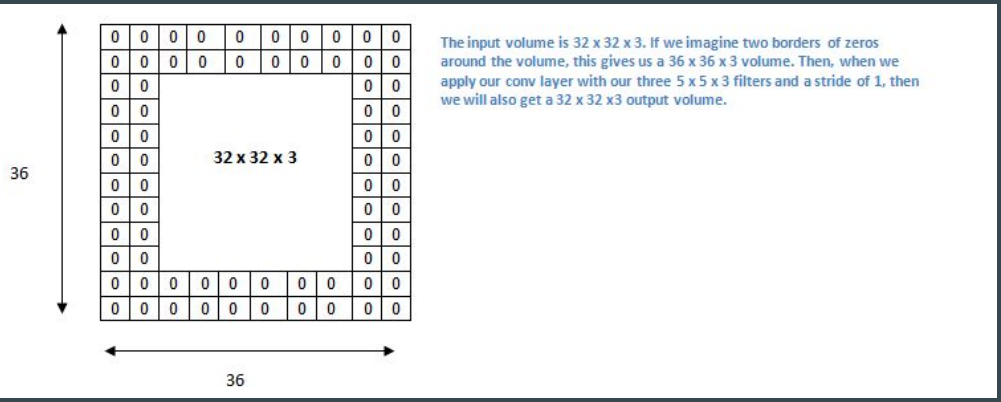

Problems (from CNN output)

Categorized Correctly

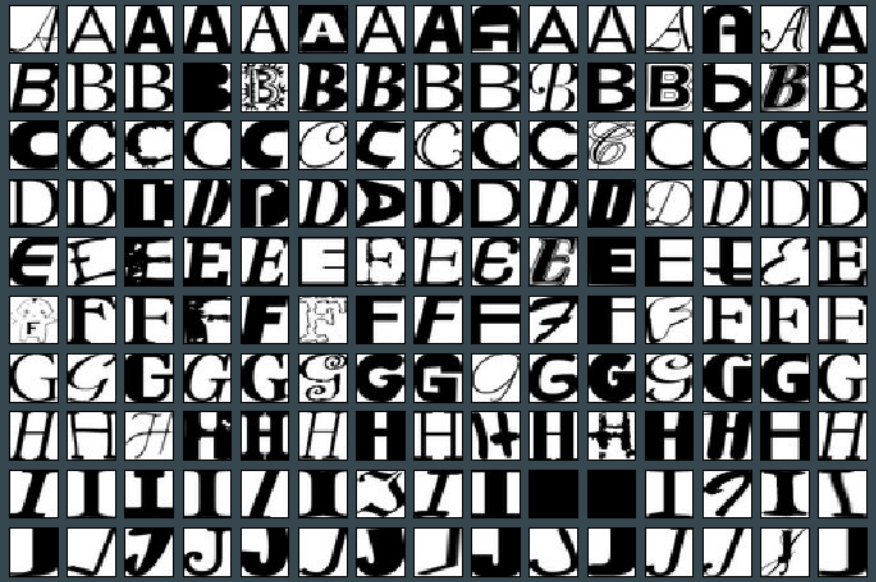

Incorrect Categorizations & Confidence Level for each Letter

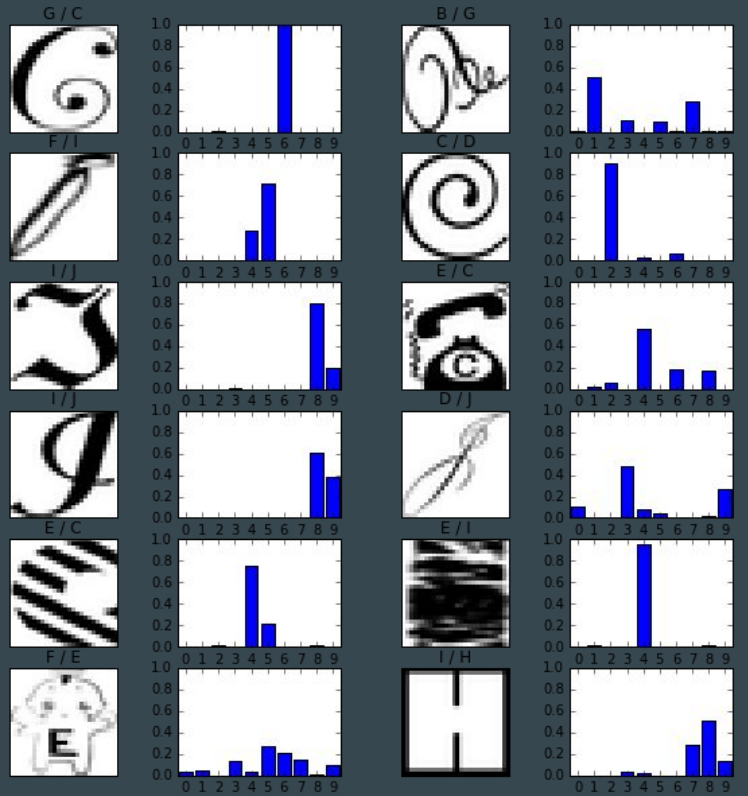

Sources

My Git Repo & Jupyter Notebook:
https://github.com/mundyreimer/notMNIST_project/blob/4f2e7b3918f585933995d6c3d227a66e6711fbb1/projects/final-project/04-notebook-rough-draft/notMNIST.ipynb

Dataset: by Yaroslav Bulatov

http://yaroslavvb.blogspot.com/2011/09/notmnist-dataset.html

Tensorflow Tutorials:
http://nbviewer.jupyter.org/github/jdwittenauer/ipython-notebooks/tree/master/notebooks/tensorflow/

Theoretical Explanation of Handwritten Digit Recognition:
https://faisalorakzai.wordpress.com/2016/06/01/handwritten-digits-recognition-using-deep-learning/

Beginner's Guide to Understanding Convolutional Neural Networks:
https://adeshpande3.github.io/adeshpande3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks/

Thanks!